

MATH 680 Computation Intensive Statistics

October 23, 2018

Gradient Descent

Contents

1	Overview of Algorithms	2
2	Line Search Methods	3
2.1	Convexity of the univariate line search problem	3
2.2	Backtracking line search	4
3	Gradient Descent	5
3.1	Why it is also called steepest descent	5
3.2	Convergence of steepest descent	7
3.3	Example: an alternative to the residual sum of squares function . . .	9

A Exact line search	19
A.1 Dichotomous search	19
A.2 Dichotomous search example: estimating the “center” of a distribution.	23
A.3 Bisection search	26
A.4 Bisection search example: estimating the “center” of a distribution continued	28
A.5 Golden section search	31

1 Overview of Algorithms

Many algorithms to compute a local or global minimizer of an objective function $f : \mathcal{X} \rightarrow \mathbb{R}$, where $\mathcal{X} \in \mathbb{R}^p$ is open, involve solving a sequence of univariate subproblems. Let $x^{(k)} \in \mathcal{X}$ be the current iterate. The next iterate is

$$x^{(k+1)} = x^{(k)} + t_k d^{(k)},$$

where $d^{(k)} \in \mathbb{R}^p$ is some direction and $t_k \in \mathbb{R}$ is some step-size. Given $d^{(k)}$, one possible choice is to use the exact line search for the step size

$$t_k = \arg \min_{t \in \mathbb{R}} g(t), \tag{1}$$

where $g(t) = f(x^{(k)} + td^{(k)})$. The feasible set in (1) could be replaced by $\{t \in \mathbb{R} : x^{(k)} + td^{(k)} \in \mathcal{X}\}$. Choosing t_k in this way ensures that $f(x^{(k+1)}) \leq f(x^{(k)})$. Approximations are used when computing (1) is difficult. The following are some

named algorithms within this framework:

- Cyclical coordinate descent uses the columns of I_p for the $d^{(k)}$'s, e.g. $d^{(1)} = (1, 0, \dots, 0)'$, $d^{(2)} = (0, 1, 0, \dots, 0)'$, etc.
- Blockwise coordinate descent. e.g. $d^{(1)} = (1, 1, 0, 0)'$, $d^{(2)} = (0, 0, 1, 1)'$, $d^{(3)} = (1, 1, 0, 0)'$, $d^{(4)} = (0, 0, 1, 1)'$.
- Steepest descent uses $d^{(k)} = -\nabla f(x^{(k)})$, and solves t_k in (1) with $t_k > 0$.
- Newton's method uses $d^{(k)} = -\{\nabla^2 f(x^{(k)})\}^{-1} \nabla f(x^{(k)})$ and $t_k = 1$.
- Quasi-Newton methods uses $d^{(k)} = -H_k^{-1} \nabla f(x^{(k)})$ where $H_k \in \mathbb{S}_+^p$ is chosen so that $d^{(k)} \approx -\{\nabla^2 f(x^{(k)})\}^{-1} \nabla f(x^{(k)})$.

2 Line Search Methods

2.1 Convexity of the univariate line search problem

Recall the univariate subproblem in 1, where the univariate objective function is defined by $g(t) = f(x^{(k)} + td^{(k)})$. If f is convex, then so is g . This is because for any two points y_1, y_2 in the domain of g and any $\lambda \in [0, 1]$,

$$\begin{aligned}
 g(\lambda y_1 + (1 - \lambda)y_2) &= f(x^{(k)} + \{\lambda y_1 + (1 - \lambda)y_2\}d^{(k)}) \\
 &= f(\lambda\{x^{(k)} + y_1 d^{(k)}\} + (1 - \lambda)\{x^{(k)} + y_2 d^{(k)}\}) \\
 &\leq \lambda f(x^{(k)} + y_1 d^{(k)}) + (1 - \lambda)f(x^{(k)} + y_2 d^{(k)}) \\
 &= \lambda g(y_1) + (1 - \lambda)g(y_2).
 \end{aligned}$$

2.2 Backtracking line search

```

backsearch <- function(f, grad.f, b, beta, alpha, quiet = FALSE,
  ...) {
  t = 1
  fk = f(b = b, ...)
  dk = grad.f(b = b, ...)
  while (1) {
    v_left = f(b = b - t * dk, ...)
    v_right = fk - alpha * t * crossprod(dk, dk)
    if (!quiet) {
      # print out condition check and current stepsize
      cat("f(x-t*dx)=", round(v_left, 2), "\n")
      cat("f(x)-at||dx||^2=", round(v_right, 2), "\n")
      cat("The current t is ", round(t, 2), "\n")
    }
    # check Armijo-Goldstein condition
    if (v_left <= v_right)
      break
    # shrink stepsize
    t = t * beta
  }
  return(t)

```

}

3 Gradient Descent

Algorithm 1 Gradient descent.

Pick an initial iterate $x^{(0)} \in \mathcal{X}$, pick a convergence tolerance $\tau > 0$, and set $k = 0$.

1. Compute $d^{(k)} = -\nabla f(x^{(k)})$.
2. If $d^{(k)} = 0$, then stop because $x^{(k)}$ is a stationary point. Otherwise, compute the step size

$$t_k = \arg \min_{t \in \mathbb{R}_+} g(t), \quad (2)$$

where $g(t) = f(x^{(k)} + td^{(k)})$.

3. Compute $x^{(k+1)} = x^{(k)} + t_k d^{(k)}$.
 4. If $\|x^{(k+1)} - x^{(k)}\| < \tau \|x^{(0)}\|$, then stop. Otherwise, replace k by $k + 1$ and go to step 1.
-

3.1 Why it is also called steepest descent

Definition 1. The directional derivatives of $f(x) = f(x_1, x_2, \dots, x_n)$ at x along a vector $d = (d_1, d_2, \dots, d_n)$ is

$$\nabla_d f(x) = \lim_{t \rightarrow 0_+} \frac{f(x + td) - f(x)}{t},$$

provided that this limit exists.

Note: If the function f is differentiable at x , then the directional derivative

exists along any vector d , and one has

$$\nabla_d f(x) = \nabla f(x)'d,$$

where ∇ on the right hand side denotes the gradient vector and $\nabla_d f(x) = \nabla f(x)'d$ is the inner product. Intuitively, the directional derivative of f at a point x represents the rate of change of f when moving past x along direction d .

Proposition 1. *Let $\mathcal{X} \subset \mathbb{R}^p$ be open. Suppose that $f : \mathcal{X} \rightarrow \mathbb{R}$ is differentiable at x and there exists a $d \in \mathbb{R}^p$ such that $\nabla_d f(x) = \nabla f(x)'d < 0$. Then for all $t > 0$ sufficiently small, $f(x + td) < f(x)$. We call this d a descent direction.*

Proof. can be easily proved by combining the definition of $\nabla_d f(x)$ and the fact that $\nabla_d f(x) = \nabla f(x)'d$ when f is differentiable. \square

Now to understand why the gradient descent algorithm is also called steepest descent, consider our current iterate $x^{(k)}$. Any direction $d \in \mathbb{R}^p$ is a descent direction if $\nabla_d f(x) = \nabla f(x^{(k)})'d < 0$ because by Proposition 1 we have $f(x^{(k)} + td) < f(x^{(k)})$ for all $t > 0$ sufficiently small. The steepest descent direction should give the smallest value in $\nabla_d f(x) = \nabla f(x^{(k)})'d$.

Suppose that $\|d\| = 1$ and $\nabla f(x^{(k)}) \neq 0$, then the unit-length steepest descent direction should return the smallest directional derivative along d , i.e.

$$\bar{d} = \arg \min_{\{d \in \mathbb{R}^p : \|d\|=1\}} \nabla f(x^{(k)})'d. \quad (3)$$

We now show that the unit-length gradient descent direction d^*

$$d^* = -\frac{\nabla f(x^{(k)})}{\|\nabla f(x^{(k)})\|}$$

is a minimizer of (3), hence $\bar{d} = d^*$. From the Cauchy–Schwartz inequality, we can obtain a lower bound for the objective function in (3): for all $d \in \mathbb{R}^p$ with $\|d\| = 1$,

$$\nabla f(x^{(k)})'d \geq -\|\nabla f(x^{(k)})\|\|d\| = -\|\nabla f(x^{(k)})\|.$$

We see that the objective function in (3) evaluated at d^* , is equal to this lower bound, so $d^* = \bar{d}$ is a global minimizer of (3). Therefore the unit-length direction of gradient descent is the unit-length direction of steepest descent.

3.2 Convergence of steepest descent

Proposition 2. *Let $\mathcal{X} \subset \mathbb{R}^p$ be open and suppose that $f : \mathcal{X} \rightarrow \mathbb{R}$ is differentiable on \mathcal{X} and ∇f is continuous on the level set $S(x_0) = \{x \in \mathcal{X} : f(x) \leq f(x_0)\}$, which we assume is closed and bounded. Let $\{x^{(k)}\}$ be the sequence of points generated by the gradient descent algorithm. Then all limit points $\{x^{(k)}\}$ are stationary points of f .*

Proof. Since $f(x^{(k+1)}) \leq f(x^{(k)})$, we have that $\{x^{(k)}\} \subset S(x_0)$, so by the Weierstrass theorem, there exists a convergent subsequence $\{x_{j(k)}\} \rightarrow \bar{x} \in S(x_0)$. Assume that $\nabla f(\bar{x}) \neq 0$. Then there exists a $\bar{t} > 0$ such that

$$\delta = f(\bar{x}) - f(\bar{x} - \bar{t}\nabla f(\bar{x})) > 0$$

and

$$\bar{x} - \bar{t}\nabla f(\bar{x}) \in \text{interior}(S(x_0)).$$

Since we assumed ∇f was continuous on $S(x_0)$,

$$\lim_{k \rightarrow \infty} \{x_{j(k)} - \bar{t}\nabla f(x_{j(k)})\} = \bar{x} - \bar{t}\nabla f(\bar{x}).$$

So for k sufficiently large,

$$\begin{aligned} f(x_{j(k)} - \bar{t}\nabla f(x_{j(k)})) &\leq f(\bar{x} - \bar{t}\nabla f(\bar{x})) + \delta/2 \\ &= f(\bar{x}) - \delta + \delta/2 \\ &= f(\bar{x}) - \delta/2, \end{aligned}$$

but

$$f(\bar{x}) \leq f(x_{j(k)} - t_{j(k)}\nabla f(x_{j(k)})) \leq f(x_{j(k)} - \bar{t}\nabla f(x_{j(k)})) \leq f(\bar{x}) - \delta/2,$$

which is impossible because $\delta > 0$. □

Remark: Gradient method does not handle non-differentiable problems. For example in Figure 1,

$$f(x) = \sqrt{x_1^2 + \gamma x_2^2} \quad (|x_2| < x_1), \quad f(x) = \frac{x_1 + \gamma|x_2|}{\sqrt{1 + \gamma}} \quad (|x_2| > x_1)$$

with exact line search, $x^{(0)} = (\gamma, 1)$, gradient method converges to non-optimal point.




Figure 1: Gradient method does not handle non-differentiable problems.

3.3 Example: an alternative to the residual sum of squares function

In a linear regression model, let $y = (y_1, \dots, y_n)'$ be the measured responses for the n cases and let $X \in \mathbb{R}^{n \times p}$ be the design matrix, where its i th row $x'_i = (1, x_{i2}, \dots, x_{ip})' \in \mathbb{R}^p$ has the values of the $p - 1$ explanatory variables for the i th case ($i = 1, \dots, n$). The model assumes that y is realization of

$$Y = X\beta_* + \epsilon,$$

where $\beta_* = (\beta_{*1}, \dots, \beta_{*p})' \in \mathbb{R}^p$ is the unknown regression coefficient vector and $\epsilon = (\epsilon_1, \dots, \epsilon_n)'$ has $\epsilon_1, \dots, \epsilon_n$ iid with an unspecified distribution having mean zero

and unknown variance σ_*^2 . Consider the estimator of β_* defined by

$$\hat{\beta}_{(\delta)} = \arg \min_{\beta \in \mathbb{R}^p} \sum_{i=1}^n |y_i - \beta' x_i|^\delta,$$

where $\delta > 1$ is user-selected. Let $f : \mathbb{R}^p \rightarrow \mathbb{R}$ be the objective function. Then

$$\nabla f(\beta) = \delta \sum_{i=1}^n |y_i - \beta' x_i|^{\delta-1} \text{sign}(\beta' x_i - y_i) x_i.$$

To use gradient descent in R to compute $\hat{\beta}_{(\delta)}$, we first define the gradient of f and the gradient of g (the univariate subproblem's objective function)

```
## define the objective function
f = function(b, y, X, delta) {
  val = sum(abs(y - X %*% b)^delta)
  return(val)
}

## define the gradient of the objective function
grad.f = function(b, y, X, delta) {
  val = delta * crossprod(X, abs(y - X %*% b)^(delta - 1) *
    sign(X %*% b - y))
  return(val)
}
```

```

## define the gradient of the univariate subproblem's
## objective function
grad.g = function(u, cur.pt, direc, ...) {
  val = sum(grad.f(cur.pt + u * direc, ...) * direc)
  return(val)
}

```

Here `grad.g` is defined for a general direction `direc`. In the steepest descent algorithm, `direc` will be $-\nabla f(\text{cur.pt})$. Let's define a function that performs steepest descent to minimize f

```

## a function that minimizes f by steepest descent
fit.delta.lm.sd = function(y, X, delta, b.start = NULL, tol = 1e-07,
  L = 1e-07, max.u = 1000, quiet = FALSE) {
  p = dim(X)[2]
  if (is.null(b.start))
    bk = rep(0, p) else bk = b.start

  k = 0
  iterating = TRUE
  while (iterating) {
    k = k + 1

    ## compute the gradient of f at our current iterate

```

```

gf.at.bk = grad.f(b = bk, y = y, X = X, delta = delta)

## check if we have converged
if (sum(abs(gf.at.bk)) < tol) {
  iterating = FALSE
} else {
  ## the direction of steepest descent is
  direc = -1 * gf.at.bk

  ## compute the best step size in this direction
  uhat = bsearch(dg = grad.g, a0 = 0, b0 = max.u, L = L,
    quiet = TRUE, cur.pt = bk, direc = direc, y = y,
    X = X, delta = delta)

  ## update our current iterate by taking this step
  bk = bk + uhat * direc
}

if (!quiet) {
  cat("k=", k, "uhat=", uhat, "gf.at.bk=", gf.at.bk,
    "\n", "\t bk=", bk, "\n")
}

}

return(list(b = bk, k = k))

```

```
}
```

Here is an example dataset

```
set.seed(680)

reps = 10000

n = 10

p = 4

sigma.star = 1/2

## create the true beta

beta.star = c(1, 0, 0, 2)

## randomly generate the design matrix

X = cbind(1, matrix(rnorm(n * (p - 1)), nrow = n, ncol = (p -
  1)))

y = X %*% beta.star + sigma.star * rnorm(n)

## ordinary least-squares estimate of beta.star

beta.hat = qr.coef(qr(x = X), y = y)

## alternative estimates of beta.star

delta = 2

quiet = TRUE
```

```

## computed with steepest descent
system.time(expr = (fitsd = fit.delta.lm.sd(y = y, X = X, delta = delta,
      L = 1e-10, quiet = quiet)))

##      user  system elapsed
##    0.029    0.000    0.029

## compare the alternative estimate to ordinary least-squares:
cbind(fitsd$b, beta.hat)

##           [,1]      [,2]
## [1,]  1.0794  1.0794
## [2,]  0.0072  0.0072
## [3,] -0.3383 -0.3383
## [4,]  1.8673  1.8673

```

Now we consider gradient descent using backtracking line search.

```

## define the objective function
f = function(b, y, X, delta) {
  val = sum(abs(y - X %*% b)^delta)
  return(val)
}

## define the gradient of the objective function

```

```

grad.f = function(b, y, X, delta) {
  val = delta * crossprod(X, abs(y - X %*% b)^(delta - 1) *
    sign(X %*% b - y))
  return(val)
}

## make a demonstrative plot for backtracking line search
bt_plot <- function(b, f = f, grad.f = grad.f, alpha, ...) {
  xi <- seq(0, 0.08, by = 0.001)
  yi <- rep(NA, length(xi))
  zi <- rep(NA, length(xi))
  ti <- rep(NA, length(xi))

  for (i in 1:length(xi)) {
    gk <- grad.f(b, y, X, delta)
    yi[i] <- f(b - xi[i] * gk, y, X, delta)
    zi[i] <- f(b, y, X, delta) - alpha * xi[i] * crossprod(gk,
      gk)
    ti[i] <- f(b, y, X, delta) - xi[i] * crossprod(gk, gk)
  }

  plot(xi, yi, type = "l")
  lines(xi, zi, lty = 2)
  lines(xi, ti, lty = 3)


```

```

}

bt_plot(b = c(2, 2, 2, 2), f = f, grad.f = grad.f, alpha = 0.5,
       y = y, X = X, delta = 2)

```



```

## a function that minimizes f by gradient descent ## using
## backtracking line search.
fit.delta.lm.bt = function(y, X, delta, b.start = NULL, tol = 1e-07,

```



```
L = 1e-07, max.u = 1000, quiet = FALSE) {  
  p = dim(X)[2]  
  if (is.null(b.start))  
    bk = rep(0, p) else bk = b.start  
  
  k = 0  
  iterating = TRUE  
  while (iterating) {  
    k = k + 1  
  
    ## compute the gradient of f at our current iterate  
    gf.at.bk = grad.f(b = bk, y = y, X = X, delta = delta)  
  
    ## check if we have converged  
    if (sum(abs(gf.at.bk)) < tol) {  
      iterating = FALSE  
    } else {  
      ## compute the best step size in this direction  
      uhat = backsearch(f = f, grad.f = grad.f, b = bk,  
        beta = 0.7, alpha = 0.5, y = y, X = X, delta = delta,  
        quiet = TRUE)  
  
      ## update our current iterate by taking this step
```

```

        bk = bk - uhat * gf.at.bk
    }

    if (!quiet) {
        cat("k=", k, "uhat=", uhat, "gf.at.bk=", gf.at.bk,
            "\n", "\t bk=", bk, "\n")
    }
}

return(list(b = bk, k = k))
}

## computed with steepest descent
system.time(expr = (fitsd1 = fit.delta.lm.bt(y = y, X = X, delta = delta,
    L = 1e-10, quiet = quiet)))

##      user  system elapsed
##  0.024    0.000    0.024

## compare the alternative estimate to ordinary least-squares:
out = cbind(fitsd1$b, fitsd$b, beta.hat)
colnames(out) <- c("Backtracking", "Steepest", "OLS")
out

##      Backtracking Steepest      OLS
## [1,]      1.0794    1.0794  1.0794
## [2,]      0.0072    0.0072  0.0072

```

## [3,]	-0.3383	-0.3383	-0.3383
## [4,]	1.8673	1.8673	1.8673

A Exact line search

We now discuss three simple algorithms for *exactly* solving the univariate subproblem in (1).

A.1 Dichotomous search

Dichotomous search algorithm minimizes a potentially non-differentiable univariate function over a closed interval.

Algorithm 2 Dichotomous search.

Suppose that the univariate function $g : \mathbb{R} \rightarrow \mathbb{R}$ to be minimized is strictly quasi-convex over $[a_0, b_0]$. Given the initial interval of uncertainty $[a_0, b_0]$, the maximum width of the final interval of uncertainty L , and the parameter $\epsilon < L/2$, set $k = 0$.

1. If $b_k - a_k < L$ then stop.
 2. Compute

$$\lambda = \frac{a_k + b_k}{2} - \epsilon, \quad \mu = \frac{a_k + b_k}{2} + \epsilon.$$
 3. If $g(\lambda) < g(\mu)$, then set $a_{k+1} = a_k$ and $b_{k+1} = \mu$.
 Otherwise, if $g(\lambda) > g(\mu)$, then set $a_{k+1} = \lambda$ and $b_{k+1} = b_k$.
 Otherwise, if $g(\lambda) = g(\mu)$, then set $a_{k+1} = \lambda$ and $b_{k+1} = \mu$.
 4. Replace k by $k + 1$ and go to step 1.
-

Figure 2 demonstrates the algorithm. After K iterations, the length of the interval of uncertainty is $0.5^K(b_0 - a_0) + 2\epsilon(1 - 0.5^K)$. It is important to chose $L > 2\epsilon$.




Figure 2: Dichotomous search

The following R code defines a function that does Dichotomous search for a minimizer of a univariate function.

```
## Dichotomous search Minimize a univariate strictly  
## quasiconvex function over the interval [a0,b0] Arguments g,  
## the function to minimize, where g(u, ...) is the function  
## evaluated at u. a0, left endpoint of the initial interval  
## of uncertainty. b0, right endpoint of the initial interval  
## of uncertainty. L, the maximum length of the final  
## interval of uncertainty. eps, search parameter, must be  
## less than L/2. quiet, should the function stay quiet?  
## ..., additional argument specifications for g Returns the  
## midpoint of the final interval of uncertainty  
dsearch = function(g, a0, b0, L = 1e-07, eps = (L/2.1), quiet = FALSE,  
  ...) {  
  mm = mean(c(a0, b0))  
  while (b0 - a0 > L) {  
    lam = mm - eps  
    mu = mm + eps  
    g.at.lam = g(lam, ...)  
    g.at.mu = g(mu, ...)  
    if (g.at.lam < g.at.mu) {  
      b0 = mu  
    } else if (g.at.lam > g.at.mu) {
```




Figure 3: Dichotomous search: case 1, case 2 and case 3.

```

    a0 = lam
} else {
    b0 = mu
    a0 = lam
}
if (!quiet)
    cat("new interval is", a0, b0, "\n")
mm = mean(c(a0, b0))
}
return(mm)
}

```

A.2 Dichotomous search example: estimating the “center” of a distribution.

Suppose that measurements of a response x_1, \dots, x_n are a realization of n independent copies of the random variable $X \sim F_X$. Consider the estimate of the center of F defined by

$$\hat{m}_\delta = \arg \min_{m \in \mathbb{R}} \sum_{i=1}^n |x_i - m|^\delta,$$

where $\delta > 0$. Let $g(m; \delta, x_1, \dots, x_n) : \mathbb{R} \rightarrow \mathbb{R}$ be the objective function, $g(m; \delta, x_1, \dots, x_n) = \sum_{i=1}^n |x_i - m|^\delta$. If $\delta \geq 1$, then $g(m; \delta, x_1, \dots, x_n)$ is convex. Also, if $\delta > 1$ then g is differentiable. We will use the dichotomous search algorithm to compute \hat{m}_δ when $\delta = 1$, $n = 9$, and $X = 10 + Z$, where Z has the t -distribution with $df = 3$.

```
## Example: Given x_1, ..., x_n, Minimize g( ; x_1, ... x_n): R ->
## R, where g(m; x_1, ..., x_n) = sum_{i=1}^n |x_i - m|^(delta).
set.seed(680)

## generate a realization of an iid sample from a heavy tailed
## distribution with mean 10.

n = 9
mu.star = 10
x.list = mu.star + rt(n, df = 3)

## Objective function to minimize
g = function(m, x.list, delta) {
```

```
len.m = length(m)
if (len.m > 1) {
  ## this case is when we want to return a vector with ith entry
  ## g(m[i], x.list)

  mat = x.list %*% t(rep(1, len.m)) - rep(1, length(x.list)) %*%
    t(m)

  val = apply(abs(mat)^delta, 2, sum)
} else {
  val = sum(abs(x.list - m)^delta)
}

return(val)
}

delta = 1

## Minimize g with Dichotomous search

mhat.1 = dsearch(g = g, a0 = min(x.list), b0 = max(x.list), quiet = TRUE,
  x.list = x.list, delta = delta)

mhat.1


## [1] 9

## this is the same as
median(x.list)
```



```
## [1] 9

## graph g
m.seq = seq(from = min(x.list), to = max(x.list), length.out = 1000)
plot(m.seq, g(m = m.seq, x.list = x.list, delta = delta), t = "l")
points(x.list, g(m = x.list, x.list = x.list, delta = delta))
```



A.3 Bisection search

Algorithm 3 Bisection search.

Given the initial interval of uncertainty $[a_0, b_0]$, and the width of the final interval of uncertainty L , set $k = 0$.

1. If $b_k - a_k < L$ then stop.

2. Compute

$$\lambda = \frac{a_k + b_k}{2}.$$

3. If $\nabla g(\lambda) > 0$, then set $a_{k+1} = a_k$ and $b_{k+1} = \lambda$.

Otherwise, if $\nabla g(\lambda) < 0$, then set $a_{k+1} = \lambda$ and $b_{k+1} = b_k$.

Otherwise, if $\nabla g(\lambda) = 0$, then stop because λ is the stationary point hence a global minimizer of g , since g is pseudo-convex.

4. Replace k by $k + 1$ and go to step 1.

After K iterations, the length of the interval of uncertainty is $0.5^K(b_0 - a_0)$. The following R code defines a function that does Bisection search for a minimizer of a univariate function.

```
## Bisection search Minimize a univariate pseduconvex function
## over the interval [a0,b0] Arguments dg, the derivative of
## function to minimize, where dg(u, ...) is this derivative
## at u. a0, left endpoint of the initial interval of
## uncertainty. b0, right endpoint of the initial interval of
## uncertainty. L, the maximum length of the final interval
## of uncertainty. quiet, should the function stay quiet?
## Returns the midpoint of the final interval of uncertainty
```

```
bsearch = function(dg, a0, b0, L = 1e-07, quiet = FALSE, ...) {  
  mm = mean(c(a0, b0))  
  ## compute gradient at the midpoint  
  while (b0 - a0 > L) {  
    dgm = dg(mm, ...)  
    if (dgm < 0) {  
      ## function is decreasing at mm new interval is [mm, b0]  
      a0 = mm  
    } else if (dgm > 0) {  
      ## function is increasing at mm new interval is [a0, mm]  
      b0 = mm  
    } else {  
      ## mm is a stationary point  
      b0 = mm  
      a0 = mm  
    }  
    if (!quiet)  
      cat("new interval is", a0, b0, "\n")  
    mm = mean(c(a0, b0))  
  }  
  return(mm)  
}
```

A.4 Bisection search example: estimating the “center” of a distribution continued

We continue the example in section A.2. Recall that we are minimizing $g(m; \delta, x_1, \dots, x_n) : \mathbb{R} \rightarrow \mathbb{R}$, where $g(m; \delta, x_1, \dots, x_n) = \sum_{i=1}^n |x_i - m|^\delta$ and we must pick a $\delta > 1$ so that g is differentiable. Using the chain rule,

$$\nabla g(m) = \delta \sum_{i=1}^n |m - x_i|^{\delta-1} \text{sign}(m - x_i).$$

We will compare the Bisection search and dichotomous search algorithms to compute \hat{m}_δ when $\delta = 1.2$ using the same `x.list` generated in section A.2.

```
## gradient of objective function to minimize
grad.g = function(m, x.list, delta) {
  len.m = length(m)
  if (len.m > 1) {
    ## this case is when we want to return a vector with ith entry
    ## grad.g(m[i], x.list, delta)
    mat = rep(1, length(x.list)) %*% t(m) - x.list %*% t(rep(1,
      len.m))
    val = delta * apply(abs(mat)^(delta - 1) * sign(mat),
      2, sum)
  } else {
    val = delta * sum(abs(m - x.list)^(delta - 1) * sign(m -
```

```
        x.list))

    }

    return(val)
}

## Minimize the objective function with Bisection search

delta = 1.2

system.time(expr = (mhat.1.2 = bsearch(dg = grad.g, a0 = min(x.list),
    b0 = max(x.list), quiet = TRUE, x.list = x.list, delta = delta)))

##      user  system elapsed
##    0.018    0.000    0.017

mhat.1.2

## [1] 9

## compare to this minimizer from Dichotomous search


system.time(expr = (mhat.1.2d = dsearch(g = g, a0 = min(x.list),
    b0 = max(x.list), quiet = TRUE, x.list = x.list, delta = delta)))

##      user  system elapsed
##    0.000    0.000    0.001

mhat.1.2d
```

```
## [1] 9

## graph grad.g
m.seq = seq(from = min(x.list), to = max(x.list), length.out = 1000)
plot(m.seq, grad.g(m = m.seq, x.list = x.list, delta = delta),
     t = "l")
points(x.list, grad.g(m = x.list, x.list = x.list, delta = delta))
```



A.5 Golden section search

Reference:

<http://ezekiel.vancouver.wsu.edu/~cs330/lectures/minimization/gold/golden-search.pdf>